

Implementing Transaction

Friday, May 31, 2019 2:15 PM

- **Scheduler** a.k.a. **Concurrency Control Manager**

- approaches
 - {
 - Locking scheduler
 - Pessimistic concurrency control
 - e.g., SQLite, SQL Server, DB2
 - Multiversion Concurrency Control (MVCC)
 - Optimistic concurrency control ← more careful
 - Postgres, Oracle: Snapshot Isolation (SI)

- Locking Scheduler (we will only discuss this in class)

- Each element has a unique lock
- Each transaction must acquire the lock to read/write that element
- If lock is taken by other transaction, wait.
- Txn must release the lock.
- By using lock we can ensure a conflict-serializable schedule.

- Action on Locks

- $L_i(x)$ Lock element x by transaction T_i .
- $U_i(x)$ Unlock element x by T_i .

- **Two phase locking (2PL)**

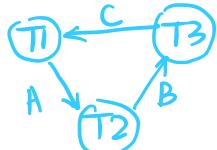
- In every transaction, every lock request must precede all unlock requests.

→ then it will be conflict serializable.

- Theorem 2PL ensures conflict serializability.

Proof: Suppose it is not serializable.

then \exists a cycle in the precedence graph.



then there's temporal cycle:

$U_1(A) \rightarrow L_2(A) \Rightarrow U_1(A)$ happened strictly before $L_2(A)$.

$L_2(A) \rightarrow U_2(B) \Rightarrow$ because 2PL.

$U_2(B) \rightarrow L_3(B)$

$L_3(B) \rightarrow U_3(C)$

$U_3(C) \rightarrow L_1(C)$

$L_1(C) \rightarrow U_1(A)$

these form a contradiction.

\therefore It must be conflict serializable.

- Non-recoverable Schedule when commit comes before rollback.

- { **Commit**: Commit the change; When a transaction is committed, it only commits its own effect.
 Rollback: roll the state back to this transaction start

- serializable

o **strict 2PL**

- All locks are held until commit/abort
- All unlocks are done together with commit/abort.

→ Strict 2PL make schedule conflict serializable and recoverable.

• **Deadlock** — another problem

e.g., $T_1: R(A), W(B)$
 $T_2: R(B), W(A)$

} holds each others lock.

- Relatively expensive: check periodically. if deadlock is found, then abort one TXN; re-check for deadlock more often.
- To find deadlock: search for a cycle.

T_1 waits for a lock held by T_2 ,

T_2 waits for a lock held by T_3 ,

...

T_n waits for a lock held by T_1 .

• **Lock Modes**

- S = shared lock (for read)

X = exclusive lock (for write)

- lock compatibility matrix

txni	tm2	N	S	X
N		✓	✓	✓
S		✓	✓	X
X		✓	X	X

because can't read/write simultaneously

• Lock Granularity

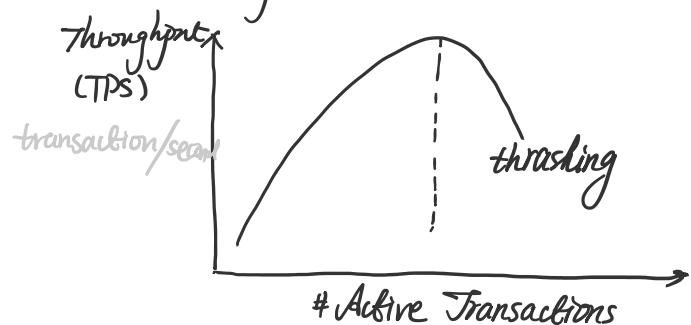
- **Fine granularity locking** (e.g. tuples)

- high concurrency
- high overhead in managing locks.
- e.g., SQL Server

- **Coarse grain locking**

- Many false conflicts
- Less overhead in managing locks
- e.g. SQLite

• Lock Performance



- **Phantom Problem** appears if tuples are inserted/deleted
 - Phantom: A tuple that is invisible during part of a transaction execution but not invisible during the entire execution.
 - Conflict-serializability assumes DB is **static**.
 - When DB is **dynamic** then c-s is not serializable.
 - **Static database**
 - Conflict serializability implies serializability.
 - **Dynamic database**
 - Conflict serializability + phantom management = serializability.
 - **(strict)2PL guarantees conflict serializability**
- Dealing with phantoms
 - Lock the entire table
 - Lock the index entry - if index is available.
 - **Predicate locks** (Complex)
 - A lock on an arbitrary predicate.
- ➔ Dealing with phantoms is expensive.